

Distributed and Fault-Tolerant System for Tuple Streaming

Jorge Santos, Miguel Vera, José Semedo
Instituto Superior Tecnico
Computer Science Department
Av. Rovisco Pais, 1, 1049-001 Lisboa
{jorge.pessoa, miguel.vera, jose.semedo}@tecnico.ulisboa.pt

Abstract

The ABSTRACT is to be in fully-justified italicized text, at the top of the left-hand column, below the author and affiliation information. Use the word “Abstract” as the title, in 12-point Times, boldface type, centered relative to the column, initially capitalized. The abstract is to be in 10-point, single-spaced type. The abstract may be up to 3 inches (7.62 cm) long. Leave two blank lines after the Abstract, then begin the main text.

1. Introduction

Data streaming distributed system architectures are used throughout both research and production environments. On one hand, for research they provide a high degree of flexibility to quickly implement and test new algorithms. On the other hand, for production they need to provide scalability (small effort of adding servers to increase performance) and high availability without the need for human intervention. Our goal was to implement a simplified fault-tolerant real-time distributed stream processing system that has strong guarantees on the correctness of its distributed computations (semantics) and the predictability of performance in case of failures (fault-tolerance).

We implemented an architecture based on tuple processing. Tuples are a data structure comprised of an arbitrary number of ordered fields (For simplicity all mentions of tuples will refer to tuples solely comprised of Strings) and a unique ID. Our tuple streaming system is comprised of a network of operators. Tuples are streamed throughout the network from a starting operator that might read them from a file to an ending operator that typically outputs to a file or produces a result. In this system a distributed computation can be seen as an acyclic graph, vertexes correspond to operators and links to data input and output. Each of these operators transforms the tuples they receive as input and may output them in a processed form. This processing

varies accordingly to the type of Operator. In our system we only support a narrow list of different operators responsible for simple operations such as filtering, counting and duplicating tuples among others. We also allow the loading of custom dll's to support the inclusion of operators tailored for specific cases.

Since this is a distributed system, operators may execute on different machines. Furthermore operators might be divided into several replicas. All replicas of a certain operator perform the same processing operations, however they might also be deployed in different machines. These replicas distribute operator load and are the core of the system's fault tolerance capabilities. Finally, the fault model we assumed when designing the system implies that only one replica will crash concurrently, as such, some of our solutions might not be adequate for systems where other types of catastrophic faults might occur.

2. Solutions

How to create a system with scaling performance, guaranteed results and highly automatized fault tolerance is not a new problem. Before we arrived at our final solution we tested and discussed other solutions for fault tolerance and semantics. Some of these solutions and their main advantages and disadvantages are briefly discussed ahead.

2.1. Waiting for an ACK

One naive approach that is often mentioned when discussing semantics and result guarantees is waiting for some sort of confirmation from the following replica confirming that the tuple has been successfully delivered to its destination. This way there is guarantee that tuples are delivered and processed correctly.

Although it is correct to say that this provides guarantees that are uncommon in distributed systems, this approach is fundamentally incorrect. Implementing this into a data

streaming distributed system would transform it into a synchronous system. This would completely annihilate(?) the purpose of a distributed data streaming architecture. The first operator in a network would have to wait until the last operator to finish for it to know that the delivery and processing were executed correctly and that it could continue processing tuples.

2.2. Mirroring data throughout the network

Another naive approach to semantics and fault tolerance is mirroring data throughout the network. In this case it could be applied by mirroring and saving received tuples on all replicas. By having all information that goes through an operator, replicas can guarantee that no information is lost when a fault occurs. Contrarily to the previous solution, this approach doesn't slow the processing of tuples by the operators, but it creates other problems.

Sending every received tuple to every other replica imposes a heavy communication load on the network infrastructure and creates the need for big storage capabilities on machines running the replicas. Despite the fact that this solution works, it is sub-optimal as keeping a copy of the tuple on every replica is unnecessary.

3. Overview

As we saw before, other solutions to this problem have several crippling disadvantages. Our analysis of other approaches allowed us to avoid some of the common pitfalls and get insight into factors like communication load imposed on the network infrastructure. Taking this into account we designed solutions we believe could be applied to real world scenarios. We tried to create an abstraction between our replica fault tolerance and our semantics algorithms, guaranteeing that one does not depend on the other, as such we feel that they should be seen as separate entities. (???????? vai contra o que o jorge diz????????)

Our Fault Tolerance algorithm can be divided into three separate phases: Fault detection, replacement and reinstating. Fault detection is handled solely by the replicas of a single operator. Replicas are organized into a ring architecture and are responsible for detecting faults on replicas placed next to them in this ring. Replica replacement is also handled by the replicas, when a replica fails, the replica that was responsible for its state monitoring completely takes over the failed replica's responsibilities. It signals the previous and following operators to ensure they no longer send data to the failed replica. Reinstating a replica is a simple process that returns the network to its original configuration. It is the opposite of the replica replacement and must again signal neighboring operators to communicate network topology changes.

Our semantics algorithm has three different behaviors for the three types of semantics: at-most-once, at-least-once, exactly-once. The at-most-once doesn't require any special behavior as in this case the system always has the same behavior irrespective of failures. At-least-once and exactly-once algorithms share a lot of their behavior and data structures. These algorithms utilize a shared tuple table on replicas to keep track of the status of all tuples received by the operator, where they were received from and where are supposed to be sent to. This table only stores the ID and status of tuples, reducing communication load in the network. Furthermore only the replicas on the path of a certain tuple store that tuple and when needed, due to the tuple table, other replicas know where it is. Finally we also implemented a confirmation system where replicas can relay a signal to previous replicas confirming the processing of the tuple, this way tuples are only stored temporarily. The main difference between exactly-once and at-least-once is that the former must guarantee that tuples are never processed twice. We utilize the tuple table to enforce that guarantee. Every time a tuple is received it is checked against the table, if it had already been processed it is not accepted.

Further and more in-depth explanation of these algorithms follows in the next chapters.

4. Fault-Tolerance

To make sure that the system is always kept in a working state, it needs to be able to recover from faults within the fault model and recover to a working network configuration. Our fault tolerance algorithm provides the ability to tolerate a downed replica of any operator, as well as proceed to its recuperation. The algorithm we used is possible due to the use of a simple, yet resourceful abstraction mechanism, which the semantics algorithm will also greatly benefit from as well. Each of the replica maintains three sets of dictionaries on which a replica is associated to a number (initially the replica number), except on the first and last operator's replicas. This is because each of the dictionaries holds the correspondence between a replica and the dictionary holder's view of it. One dictionary abstracts all of the colleague replicas of the holder (replicas of the same operator as itself), another distinct dictionary abstracts the replicas of the above operator, and again for the one below. So every replica of every operator has the ability to establish a connection with what they think is every replica above and below them, as well as their colleagues.

4.1. Fault Detection

Going into more detail on how a downed replica can be detected. Colleague replicas among themselves have the illusion of being in a ring like topology. Each replica periodi-




Figure 1. How colleague replicas ping each other

cally pings the following colleague like illustrated in Figure. 1. A replica can efficiently know what replica it precedes. It knows it's own replica index (i) therefore it knows the index of the replica it has to ping: $\text{ping}(i)$, by applying the following rule:

$$\text{ping}(i) = (i + 1) \% \text{number of replicas}$$

If the pinged replica is detected to be down, then the replica can just calculate the following replica's index using the very same rule, and contact it, asking for it to take over the downed replica's spot.

4.2. Take Over Procedure

The replica taking over goes through the 3 sets of replica holding abstractions aforementioned. Firstly it iterates over the the replica holding abstractions corresponding to the operator above it, and for each one, it substitutes the place of the replica detected as downed with a representation of itself. It literally takes over it's spot, so every command that would be aimed at the downed replica, is routed to the one that took over it's functions. It then repeats the process for it's colleague replicas and the ones of the operator below it.

There is however a small detail that caused a somewhat meaningful problem. In the case of an operator having only two replicas, and one of them going down, the only remaining replica would substitute the representation of it's colleague with itself, and then effectively start pinging itself. This caused not only some performance issues besides other problems. In order to fix this efficiently we had to break a small layer of abstraction and add another structure to the replicas, a list containing the indexes of the replicas it has taken over. This allows a replica to know whether it would be attempting to connect to itself or not.

After a replica has fully taken over for another one, all other replicas have the illusion of the topology state remaining the same. This is because every method that implies connecting to another replica, obtains the connection through the abstraction layer previously described, so every

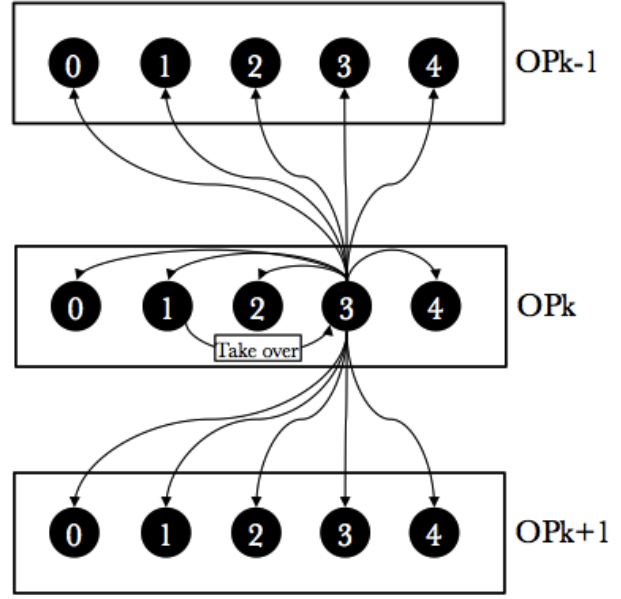


Figure 2. Take over procedure illustrated

work load is seamlessly diverted to the new replica. Only replicas within the same operator can detect downed replicas. If any other remoting method fails due to the target replica being down, it enters an active wait until another replica takes over and the method is successfully carried out. We chose this implementation over the option of any remoting task possibly triggering a take over, due to the possibility of the retrying of the faulting method being faster than the take over process. Fixing this issue would over complicate things, and the simpler implementation as we have now revealed it self as the faster one.

4.3. Reinstatement Procedure

The reinstating protocol is simply put, the reverse of the take over process. Once a replica is brought back it goes through the same process of accessing previous, colleague and following operator's replica abstraction holding structures in that respective order, and inserting a representation of itself in it's rightful spot. After this, all remote methods are once again routed correctly to the reinstated replica. The replica that had taken over for it also has it's index list updated, meaning that the index of the reinstated replica is removed from that list.

5. Semantics

One of the biggest concerns in distributed tuple processing is being able to guarantee certain semantics associated

with the processing of a tuple in the presence of system failures, and inevitably the way it needs to reconfigure in order to continue functioning properly. This issue is inevitably linked to the way that the system does fault tolerance. The algorithm used for any semantic guarantee is very similar, using the same data structures and most of the same procedures. We will first explain the general idea behind the algorithm, then used data structures and the algorithm and finally how it used to assure that a tuple is processed at most once, at least once and exactly once.

5.1. Explanation

If the system must process a certain tuple at least once then there must a guarantee that in any kind of a node failure the tuples that weren't successfully sent to a node in failure are sent either to another node or to the the original node if it recovers. There are however other scenarios in which a node can fail, For example after receiving a tuple but before sending it to the next operator. Keeping in mind that the system must obviously be asynchronous in this confirmation, then the previous node can't easily know it has to resend the tuple, and when it doesn't, or how long it needs to keep the tuple. In this approach the tuples would need to be kept in every operator at least until the tuple was processed by every node, presenting scalability issues. That's why we opted to use an algorithm based in keeping the information of the tuples being processed and by which replica, as well as which replica originally contains the tuple. That way all the tuples that weren't sent can be processed and sent again if a new node takes over a dead node without the need to replica tuples by all nodes of an operator.

5.2. Data Structures

To understand how the algorithm works we must first explain the used data structures: A tuple is identified by it's Tuple Id structure which represents a stream of a tuple along the processing chain. In each specific tuple there is an unique id that is usually kept along all operators, unless there must an output of several tuples from the same one, efectively diverging the tuple stream. The remaining information kept in the Tuple Id refers to the operator and replica it came from.

In each node there is a delivery table (a simple Hash-based Map that stores each tuple by its unique id) and that keeps every tuple received in a replica until it can be disposed.

There is also a shared tuple table in each node that stores Tuple Records associated to a tuple id (in a simillar Hash Map). A tuple record is a small representation of a tuple containing its Tuple Id, as well as the replica emitting the tuple record. This way the tuple record contains only the

necessary information for a node to re-process a tuple that wasn't properly processed due to failure. For information storing purposes the tuple records have a state (pending or purged) that is used to store tuple records of the tuples processed by a replica in a purged state instead of deleting them in order to properly know what tuples have already been processed by that replica.

5.3. Algorithm




Figure 3. Ordered steps taken when a node processes a tuple

The ordered steps represented in figure 1 show most of the necessary procedures executed every time a tuple is received, to keep the data structures synchronized in order to allow a proper recovery when a node fails. When a node receives a tuple (1) the first thing it must do is insert it into it's delivery table (2) and it's shared tuple record table (3). The node then synchronizes every table with other replicas shared tables (4) and finally processes the tuple, sending it afterwards (5). After the tuple is sent and confirmation is received (6) the node will first issue the deletion of the tuple record of every other shared table and purge its own shared table of the same tuple (7). Finally it will warn the replica which the tuple originated from that it can finally delete it from its delivery table (8) since it was already received by the forward node. This information is used whenever a node takes over another failing node. It scans its shared table looking for tuple records of the dead node and for each record it finds, it requests the tuple to the replica that the tuple originated from. Finally that tuple is processed and sent to the following node, by the same procedures explained before, as if a new tuple was received.

Note that this approach only ensures the semantics of the delivery in case of a single horizontal failure (when two sequential replicas on the path of processing a tuple fail). It

could also support more failures if necessary by extending the number of nodes where a tuple is kept before it is delivered.

5.4. Differences for each guarantee

5.4.1 At-Most-Once

To assure that a tuple is processed at most once, a system must only send any tuple once, independent of failure. Since a configuration for this system is dependent on being an acyclic graph, the only guarantee we need to provide is that there isn't any kind of mechanism to resend tuples in case of a takeover of a node that crashed by another node. As such the implementation of this strategy relies on not sharing any kind of information about tuples already processed, which is achieved by not executing any kind of the procedures explained before related to the applied algorithm (that isn't really applied in this case of semantics). This way, a certain tuple from an input is only sent once in the forward direction of the distributed network, and in case of a node failing all the tuples that it had processed but not sent yet won't be processed at all. The same applies to a replica that sends a tuple to the next replica, which in case of not getting a response, will just drop the tuple.

5.4.2 At-Least-Once

The algorithm previously basically accomplishes at least once delivery by itself, without the need of any changes.

5.4.3 Exactly-Once

Finally to guarantee that tuples are delivered only exactly once, and taking the algorithm described as a starting point, there is the need to guarantee that a tuple is never processed twice on the same operator. The first thing to remember is that each tuple has a unique Id that identifies its path along the processing chain. Based on that, each received tuple on a node is checked by its id to guarantee that it has never been processed, by comparing its unique id to every tuple record stored in the shared tuple record table (hence the possible states for a certain tuple). The tuple is accepted with success in case it was never inserted in the table, assuring the replica that the tuple was never processed and sent to next operator before, which together with the guarantees provided by the algorithm assure exactly once delivery of tuples.

6. Evaluation

In order to properly understand the effectiveness of the developed solution it is necessary to understand the performance implications of the used algorithm and how it com-

pares with other possible solutions. Firstly we will analyze the used solution from a theoretical standpoint and compare it with the previously mentioned naive solution. Then we will test a few measures and compare them to the theoretical results in order to provide some confirmation.

6.1. Analysis

To simplify the process of creating a model that fits the developed system we will consider a few assumptions that hopefully won't distort the obtained results by much:

- The considered system is only comprised of simple operators where one tuple can only originate another tuple, without stream divergences;
- The system can be simplified by some measures where R is the average replication factor of each operator, O is the number of operators and T is the average size of a tuple being sent;
- Tuple Records have a constant size of 64 bits for the unique id plus 32 bits to identify the operator and replica and 1 bit for the state. This value is represented by $TR \approx 128 \text{ bits}$.

In this case the system consists of an acyclic (forward) graph where we can attribute a weight to each transition depending on what measure we're analyzing.

Let's start by considering the amount of data transferred per tuple sent: For each tuple received by a replica it sends a tuple record to every other replica, so $(R-1) \times TR$. There is also a purge message that contains only a unique id, which we will approximate to the Tuple Record (TR) size. Similarly the confirmation to the previous contains the id. All things considered the total amount of data for each tuple received per replica for all the operators is

$$(R - 1 + 1 + 1) \times TR \times O \Rightarrow O(R \times TR \times O)$$

$$TR \text{ is constant} \Rightarrow O(R \times O)$$

Comparatively a naive solution that receives a tuple on a replica and replicates it through all the replicas (removing the necessity for forward and backwards confirmation) would require

$$(R - 1) \times T \times O \Rightarrow O(R \times T \times O)$$

The difference of both approaches being that our proposed solution scales with the size of a tuple record, TR which is a constant value, while the naive approach will scale with the tuple size as well T . For increasing tuple sizes the difference in data that needs to be exchanged between replicas for each tuple is noticeably smaller as we expect to confirm in the practical results. In this sense, we can start to understand the performance advantages presented by our solution when $T \gg TR$.

6.2. Qualitative

Qualitative evaluation

7. Discussion

Prior to analyzing the used algorithms, we have to acknowledge that the solutions aren't perfect but are a step forward from the naive approaches that might be picked in this situation. Most of the differences are actually unintelligible when processing small datasets, in which each tuple is small.

For example in this case there is an advantage of replicating tuples in each node after receiving from the operator, considering their sizes might be even smaller than the tuple records and there are less steps involved in the algorithm. But on the other hand, if we try to scale the tuple size even a little bit, then the tuple-replicating approach would scale much worse, resulting in huge amounts of exchanged data between replicas of the same operator before they can even process the tuples. On the other hand, replicating only tuple records scales in a constant way with tuple size (doesn't increase). So as analyzed from this feature, one thing to note is that there is no perfect solution that applies to every kind of system in a scalable and efficient way. A much better compromise is to develop the system according to whatever are its needs, and if possible without forgetting the scalability to a real word scenario of applying the developed solution.

As far as fault detection goes, we think that other topologies can be adopted besides our ring-like one. It would really come down to the real life application of the system. A possibly good heuristic could be how far replicas are from each other, so the whole operator ping cycle (all replicas get pinged) is made as fast as possible. With different topologies different methods of picking what replica takes over are necessary. If the topology is a tree-like, then choosing a candidate becomes much harder. Some other things can be taken into account when picking a take over candidate, for example, work load metrics can be shared with each "are you alive" in order to pick which replica would take the extra load better. Another even more complex possible solution is to have more than one replica assigned to taking over the downed replica's work load. Once again it all depends on the real life application scenario. Regarding the take-over and reinstating methods, we support the opinion that there aren't any meaningful variations of our adopted solution that can be done taking into account our network topology.

8. Conclusion

Evaluation of the solution